

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
import plotly.express as px
from sklearn.linear_model import LogisticRegression
from sklearn.cluster import KMeans
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import LabelEncoder
import warnings
warnings.filterwarnings("ignore")

```

2. Load the dataset into the tool.

```

from google.colab import files
upload=files.upload()

```

<IPython.core.display.HTML object>

Saving Mall_Customers.csv to Mall_Customers.csv

```
mydata = pd.read_csv('Mall_Customers.csv')
```

```
mydata.head()
```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	Male	19	15	39
1	2	Male	21	15	81
2	3	Female	20	16	6
3	4	Female	23	16	77
4	5	Female	31	17	40

```
mydata.tail()
```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
195	196	Female	35	120	
196	197	Female	45	126	
197	198	Male	32	126	
198	199	Male	32	137	
199	200	Male	30	137	

```
mydata.describe()
```

	CustomerID	Age	Annual Income (k\$)	Spending Score (1-100)
count	200.000000	200.000000	200.000000	

```

200.000000
mean    100.500000    38.850000    60.560000
50.200000
std      57.879185    13.969007    26.264721
25.823522
min      1.000000    18.000000    15.000000
1.000000
25%     50.750000    28.750000    41.500000
34.750000
50%     100.500000    36.000000    61.500000
50.000000
75%     150.250000    49.000000    78.000000
73.000000
max     200.000000    70.000000   137.000000
99.000000

```

```
mydata.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 5 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   CustomerID            200 non-null    int64
 1   Gender                200 non-null    object
 2   Age                  200 non-null    int64
 3   Annual Income (k$)    200 non-null    int64
 4   Spending Score (1-100) 200 non-null    int64
dtypes: int64(4), object(1)
memory usage: 7.9+ KB

```

```
mydata.dtypes
```

```

CustomerID            int64
Gender                object
Age                  int64
Annual Income (k$)    int64
Spending Score (1-100) int64
dtype: object

```

3.Perform visualisations Univariate Analysis





```
mydata.hist(figsize=(20,10), grid = False, bins = 50)
```

```


array([[<matplotlib.axes._subplots.AxesSubplot object at
0x7f9dff9e9450>,
      <matplotlib.axes._subplots.AxesSubplot object at
0x7f9dff9e0610>],
      [<matplotlib.axes._subplots.AxesSubplot object at
0x7f9dff999c10>,
      <matplotlib.axes._subplots.AxesSubplot object at



```

```
0x7f9dff95d250>]],  
dtype=object)
```






```
columns = ['Age', 'Annual Income (k$)', 'Spending Score (1-100)']  
for i in columns:  
    plt.figure()  
    sns.distplot(mydata[i])
```








```
columns = ['Age', 'Annual Income (k$)', 'Spending Score (1-100)']
for i in columns:
    plt.figure()
    sns.kdeplot(mydata[i], shade=True, hue=mydata['Gender'])
```





```
columns = ['Age', 'Annual Income (k$)', 'Spending Score (1-100)']
for i in columns:
    plt.figure()
    sns.boxplot(data=mydata, x='Gender', y=mydata[i])
```





```
mydata['Gender'].value_counts(normalize=True)
```

```
Female    0.56
```

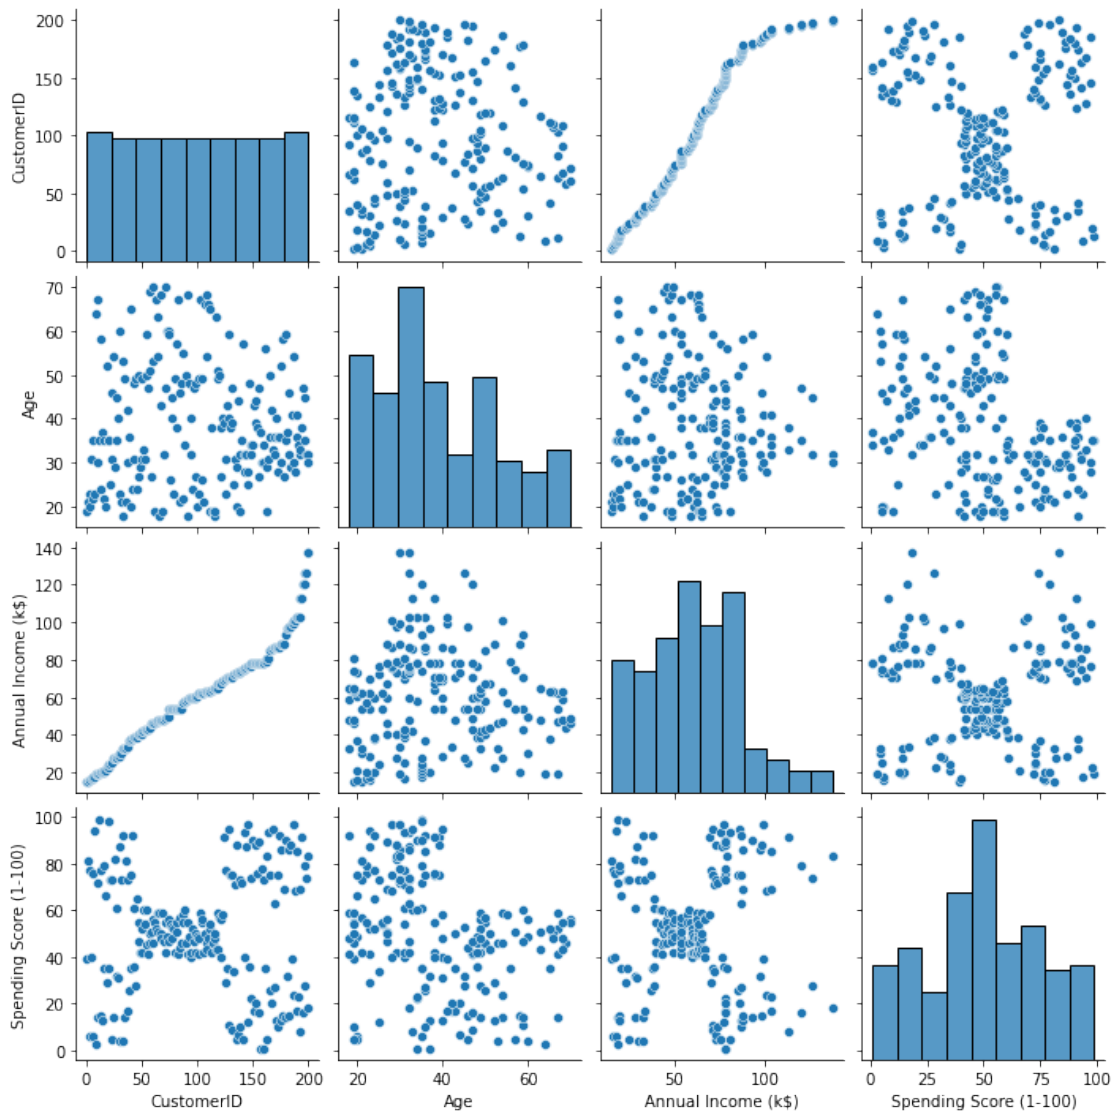
```
Male      0.44
```

```
Name: Gender, dtype: float64
```

Bi- Variate Analysis


```
numerical_features = mydata.select_dtypes(include=[np.number]).columns
categorical_features =
mydata.select_dtypes(include=[np.object]).columns
numerical_features
categorical_features
sns.pairplot(mydata[numerical_features])
```

<seaborn.axisgrid.PairGrid at 0x7f9dfedfe090>



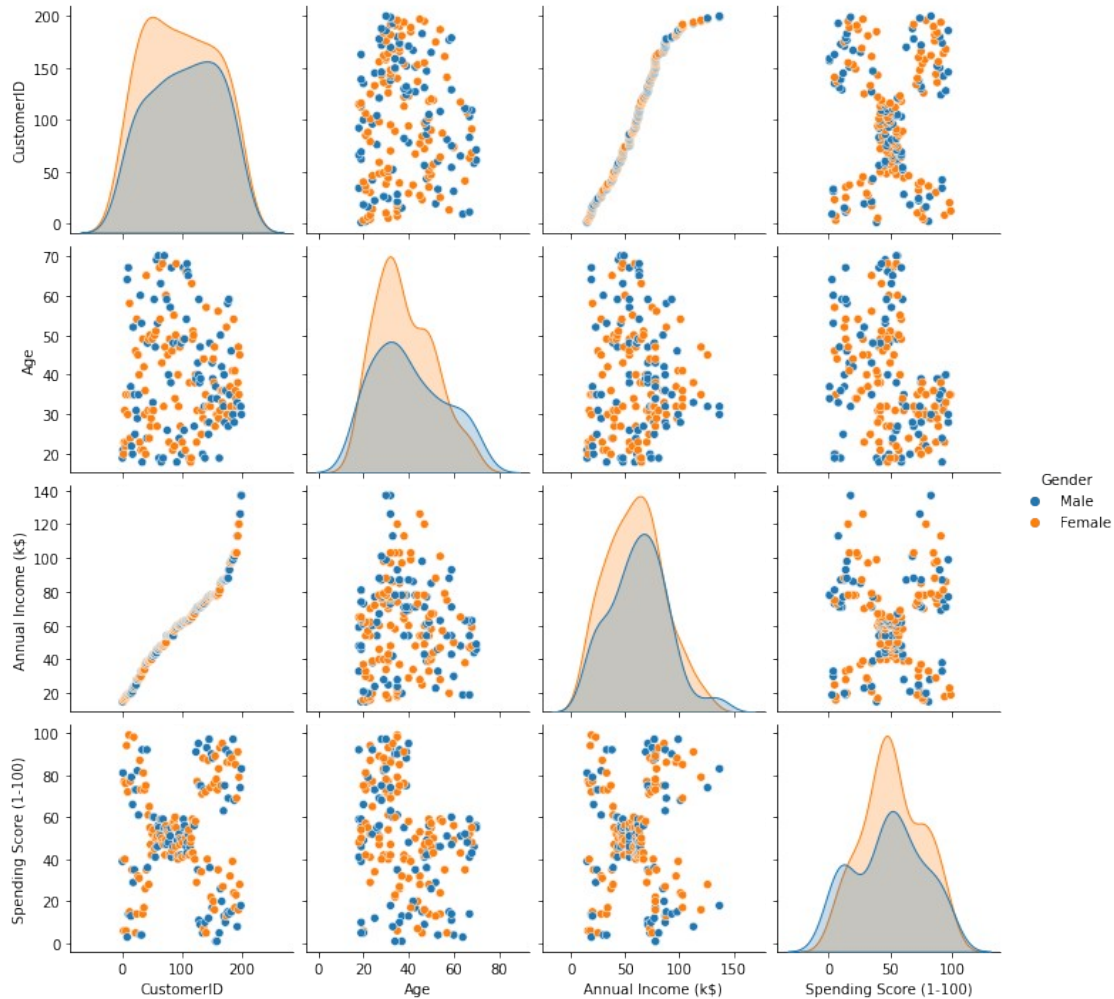
```
sns.scatterplot(data=mydata, x='Annual Income (k$)', y='Spending Score (1-100)')
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f9dff837910>



```
sns.pairplot(mydata,hue='Gender')
```

```
<seaborn.axisgrid.PairGrid at 0x7f9dff8d1890>
```



```
mydata.groupby(['Gender'])['Age', 'Annual Income (k$)', 'Spending Score (1-100)'].mean()
```

	Age	Annual Income (k\$)	Spending Score (1-100)
Gender			
Female	38.098214	59.250000	51.526786
Male	39.806818	62.227273	48.511364

```
mydata.columns
```

```
Index(['CustomerID', 'Gender', 'Age', 'Annual Income (k$)',  
      'Spending Score (1-100)'],  
      dtype='object')
```

```
mydata.corr()
```

	CustomerID	Age	Annual Income (k\$)	\
CustomerID	1.000000	-0.026763	0.977548	
Age	-0.026763	1.000000	-0.012398	
Annual Income (k\$)	0.977548	-0.012398	1.000000	
Spending Score (1-100)	0.013835	-0.327227	0.009903	

```

                                Spending Score (1-100)
CustomerID                      0.013835
Age                             -0.327227
Annual Income (k$)              0.009903
Spending Score (1-100)         1.000000

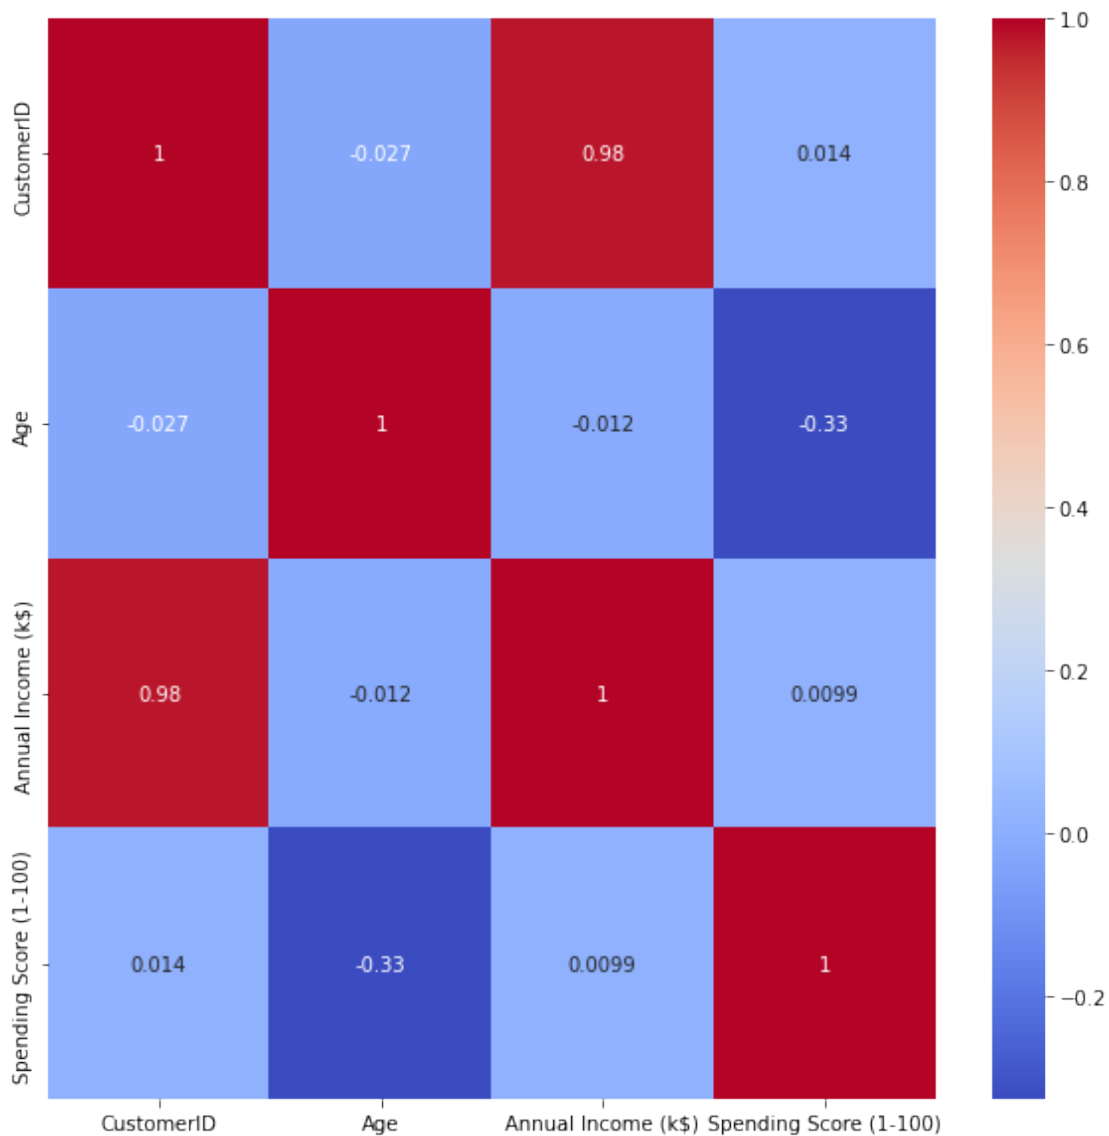
```

Multivariate Analysis

```

plt.figure(figsize=(10, 10))
sns.heatmap(mydata.corr(),annot=True,cmap='coolwarm')
<matplotlib.axes._subplots.AxesSubplot at 0x7f9dfedc3f90>

```



4.Perform descriptive statistics on the dataset

```
mydata.mean()
```

```
CustomerID          100.50
Age                  38.85
Annual Income (k$)   60.56
Spending Score (1-100) 50.20
dtype: float64
```

```
mydata.std()
```

```
CustomerID          57.879185
Age                  13.969007
Annual Income (k$)   26.264721
Spending Score (1-100) 25.823522
dtype: float64
```

```
mydata.sum()
```

```
CustomerID
20100
Gender
MaleMaleFemaleFemaleFemaleFemaleFemaleFemaleMale...
Age
7770
Annual Income (k$)
12112
Spending Score (1-100)
10040
dtype: object
```

```
mydata.count()
```

```
CustomerID          200
Gender              200
Age                200
Annual Income (k$)  200
Spending Score (1-100) 200
dtype: int64
```

```
mydata.min()
```

```
CustomerID          1
Gender              Female
Age                18
Annual Income (k$)  15
Spending Score (1-100) 1
dtype: object
```

```
mydata.var()
```

```
CustomerID          3350.000000
Age                195.133166
Annual Income (k$)  689.835578
```

```
Spending Score (1-100)      666.854271
dtype: float64
```

```
mydata.mode()
```

```
      CustomerID  Gender  Age  Annual Income (k$)  Spending Score (1-
100)
0             1  Female  32.0                54.0
42.0
1             2     NaN   NaN                78.0
NaN
2             3     NaN   NaN                NaN
NaN
3             4     NaN   NaN                NaN
NaN
4             5     NaN   NaN                NaN
NaN
..           ...     ...   ...                ...
...
195          196     NaN   NaN                NaN
NaN
196          197     NaN   NaN                NaN
NaN
197          198     NaN   NaN                NaN
NaN
198          199     NaN   NaN                NaN
NaN
199          200     NaN   NaN                NaN
NaN
```

```
[200 rows x 5 columns]
```

```
mydata.median()
```

```
CustomerID      100.5
Age             36.0
Annual Income (k$)  61.5
Spending Score (1-100)  50.0
dtype: float64
```

5. Check for Missing values and deal with them.

```
mydata.duplicated().sum()
```

```
0
```

```
mydata.isna().sum()
```

```
CustomerID      0
Gender          0
Age             0
Annual Income (k$)  0
```

```
Spending Score (1-100)    0
dtype: int64
```

```
mydata.nunique()
```

```
CustomerID                200
Gender                    2
Age                      51
Annual Income (k$)        64
Spending Score (1-100)    84
dtype: int64
```

```
mydata.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 200 entries, 0 to 199
```

```
Data columns (total 5 columns):
```

#	Column	Non-Null Count	Dtype
0	CustomerID	200 non-null	int64
1	Gender	200 non-null	object
2	Age	200 non-null	int64
3	Annual Income (k\$)	200 non-null	int64
4	Spending Score (1-100)	200 non-null	int64

```
dtypes: int64(4), object(1)
```

```
memory usage: 7.9+ KB
```

```
mydata.sum()
```

```
CustomerID
20100
Gender
MaleMaleFemaleFemaleFemaleFemaleFemaleFemaleMa...
Age
7770
Annual Income (k$)
12112
Spending Score (1-100)
10040
dtype: object
```

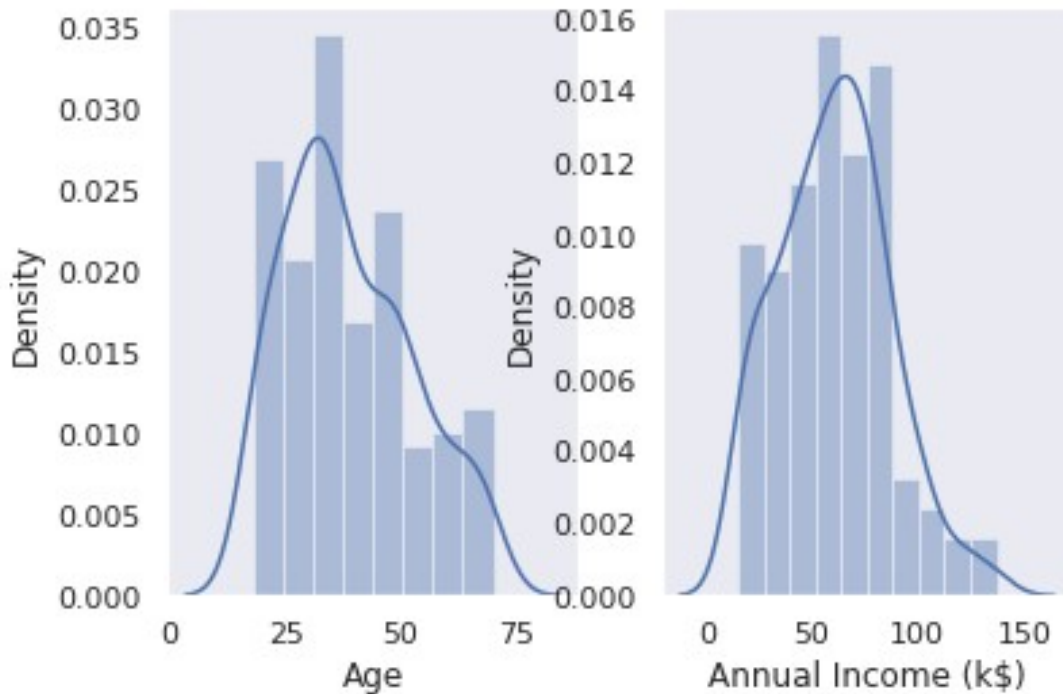
```
mydata.isnull().any()
```

```
CustomerID    False
Gender        False
Age           False
Annual Income (k$)  False
Spending Score (1-100) False
dtype: bool
```

6. Find the outliers and replace them outliers

```
sns.set(style = 'dark')
plt.subplot(1,2,1)
sns.distplot(mydata['Age'])
plt.subplot(1,2,2)
sns.distplot(mydata['Annual Income (k$)'])
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f9dff5ab810>



```
print("Highest allowed",mydata['Age'].mean() + 3*mydata['Age'].std())
print("Lowest allowed",mydata['Age'].mean() - 3*mydata['Age'].std())
```

Highest allowed 80.75702199467665
Lowest allowed -3.0570219946766386

```
new_mydata = mydata[(mydata['Age'] < 50) & (mydata['Age'] > 25)]
new_mydata
```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
4	5	Female	31	17	
40					
6	7	Female	35	18	
6					
9	10	Female	30	19	
72					
11	12	Female	35	19	
99					
14	15	Male	37	20	
13					

```

..          ...      ...      ...          .
..
195          196   Female    35          120
79
196          197   Female    45          126
28
197          198     Male    32          126
74
198          199     Male    32          137
18
199          200     Male    30          137
83

```

```
[117 rows x 5 columns]
```

```

upper_limit = mydata['Age'].mean() + 5*mydata['Age'].std()
lower_limit = mydata['Age'].mean() - 5*mydata['Age'].std()

```

```

mydata['Age'] = np.where(
    mydata['Age']>upper_limit,
    upper_limit,
    np.where(
        mydata['Age']<lower_limit,
        lower_limit,
        mydata['Age']
    )
)

```

```
mydata['Age'].describe()
```

```

count    200.000000
mean      38.850000
std       13.969007
min       18.000000
25%       28.750000
50%       36.000000
75%       49.000000
max       70.000000
Name: Age, dtype: float64

```

7. Check for Categorical columns and perform encoding.

```
mydata.columns
```

```

Index(['CustomerID', 'Gender', 'Age', 'Annual Income (k$)',
      'Spending Score (1-100)'],
      dtype='object')

```

```

mydata_numeric = mydata[['CustomerID', 'Age', 'Annual Income
(k$)', 'Spending Score (1-100)']]
mydata_categorical = mydata[['Gender']]

```



```
mydata_numeric.head()
```

	CustomerID	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	19.0	15	39
1	2	21.0	15	81
2	3	20.0	16	6
3	4	23.0	16	77
4	5	31.0	17	40

```
mydata_categorical.head()
```

	Gender
0	Male
1	Male
2	Female
3	Female
4	Female

```
print(mydata['CustomerID'].unique())
print(mydata['Gender'].unique())
print(mydata['Age'].unique())
print(mydata['Annual Income (k$)'].unique())
print(mydata['Spending Score (1-100)'].unique())
```

```
[ 1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17
18
19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35
36
37 38 39 40 41 42 43 44 45 46 47 48 49 50 51 52 53
54
55 56 57 58 59 60 61 62 63 64 65 66 67 68 69 70 71
72
73 74 75 76 77 78 79 80 81 82 83 84 85 86 87 88 89
90
91 92 93 94 95 96 97 98 99 100 101 102 103 104 105 106 107
108
109 110 111 112 113 114 115 116 117 118 119 120 121 122 123 124 125
126
127 128 129 130 131 132 133 134 135 136 137 138 139 140 141 142 143
144
145 146 147 148 149 150 151 152 153 154 155 156 157 158 159 160 161
162
163 164 165 166 167 168 169 170 171 172 173 174 175 176 177 178 179
180
181 182 183 184 185 186 187 188 189 190 191 192 193 194 195 196 197
198
199 200]
['Male' 'Female']
[19. 21. 20. 23. 31. 22. 35. 64. 30. 67. 58. 24. 37. 52. 25. 46. 54.
29.
45. 40. 60. 53. 18. 49. 42. 36. 65. 48. 50. 27. 33. 59. 47. 51. 69.]
```

```

70.
63. 43. 68. 32. 26. 57. 38. 55. 34. 66. 39. 44. 28. 56. 41.]
[ 15 16 17 18 19 20 21 23 24 25 28 29 30 33 34 37 38
39
40 42 43 44 46 47 48 49 50 54 57 58 59 60 61 62 63
64
65 67 69 70 71 72 73 74 75 76 77 78 79 81 85 86 87
88
93 97 98 99 101 103 113 120 126 137]
[39 81 6 77 40 76 94 3 72 14 99 15 13 79 35 66 29 98 73 5 82 32 61
31
87 4 92 17 26 75 36 28 65 55 47 42 52 60 54 45 41 50 46 51 56 59 48
49
53 44 57 58 43 91 95 11 9 34 71 88 7 10 93 12 97 74 22 90 20 16 89
1
78 83 27 63 86 69 24 68 85 23 8 18]

```

```
gender_encoder = LabelEncoder()
```

```
gender_encoder.fit(mydata_categorical['Gender'])
```

```
LabelEncoder()
```

```
gender_values = gender_encoder.transform(mydata_categorical['Gender'])
```

```
print("Before Encoding:", list(mydata_categorical['Gender'][-10:]))
```

```
print("After Encoding:", gender_values[-10:])
```

```
print("The inverse from the encoding result:",
```

```
gender_encoder.inverse_transform(gender_values[-10:]))
```

```
Before Encoding: ['Female', 'Female', 'Male', 'Female', 'Female',
'Female', 'Female', 'Male', 'Male', 'Male']
```

```
After Encoding: [0 0 1 0 0 0 0 1 1 1]
```

```
The inverse from the encoding result: ['Female' 'Female' 'Male'
```

```
'Female' 'Female' 'Female' 'Female' 'Male'
```

```
'Male' 'Male']
```

8. Scaling the data

```
numCol = [col for col in mydata.columns if mydata[col].dtype != "0"]
```

```
numCol
```

```
['CustomerID', 'Age', 'Annual Income (k$)', 'Spending Score (1-100)']
```

```
catColumn = [col for col in mydata.columns if mydata[col].dtype ==
"0"]
```

```
catColumn
```

```
['Gender']
```

```

from sklearn.preprocessing import scale
X=mydata.drop(columns=['Age'],axis=1)
X.head()

```

	CustomerID	Gender	Annual Income (k\$)	Spending Score (1-100)
0	1	Male	15	39
1	2	Male	15	81
2	3	Female	16	6
3	4	Female	16	77
4	5	Female	17	40

```

scaler=MinMaxScaler()

```

```

mydata[["Age"]]=scaler.fit_transform(mydata[["Age"]])
print(mydata)

```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	Male	0.019231	15	39
1	2	Male	0.057692	15	81
2	3	Female	0.038462	16	6
3	4	Female	0.096154	16	77
4	5	Female	0.250000	17	40
..
195	196	Female	0.326923	120	79
196	197	Female	0.519231	126	28
197	198	Male	0.269231	126	74
198	199	Male	0.269231	137	18
199	200	Male	0.230769	137	83

```

[200 rows x 5 columns]

```

9. Perform any of the clustering algorithms

```

clustering1 = KMeans(n_clusters=6)
clustering1.fit(mydata[['Annual Income (k$)']])
KMeans(n_clusters=6)
clustering1.labels_

```

```

array([3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3,
3,
      3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1,
      1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1,
      1, 1, 1, 1, 1, 1, 1, 1, 1, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4,
4,
      4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4,
4,
      4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5,
5,
      5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 2, 2, 2, 2, 2,
2,
      2, 2], dtype=int32)

```

```

mydata['Income Cluster']=clustering1.labels_
mydata.head()

```

	CustomerID (1-100) \	Gender	Age	Annual Income (k\$)	Spending Score
0	1	Male	0.019231	15	
39					
1	2	Male	0.057692	15	
81					
2	3	Female	0.038462	16	
6					
3	4	Female	0.096154	16	
77					
4	5	Female	0.250000	17	
40					

	Income Cluster
0	3
1	3
2	3
3	3
4	3

```

mydata['Income Cluster'].value_counts()

```

4	48
1	42
0	42
3	32
5	28

```
2      8
Name: Income Cluster, dtype: int64
```

```
clustering1.inertia_
```

```
5050.9047619047615
```

10. Add the cluster data with the primary dataset

```
inertia_scores=[mydata]
for i in range(1,11):
    kmeans=KMeans(n_clusters=i)
    kmeans.fit(mydata[['Annual Income (k$)']])
    inertia_scores.append(kmeans.inertia_)
```

```
inertia_scores
```

```
[      CustomerID  Gender      Age  Annual Income (k$)  Spending Score
(1-100) \
0              1    Male  0.019231              15
39
1              2    Male  0.057692              15
81
2              3  Female  0.038462              16
6
3              4  Female  0.096154              16
77
4              5  Female  0.250000              17
40
..           ...      ...      ...           ...
...
195           196  Female  0.326923              120
79
196           197  Female  0.519231              126
28
197           198    Male  0.269231              126
74
198           199    Male  0.269231              137
18
199           200    Male  0.230769              137
83
```


```
      Income Cluster
0              3
1              3
2              3
3              3
4              3
..           ...
195           2
196           2
```

```
197          2
198          2
199          2
```

```
[200 rows x 6 columns],
137277.280000000006,
48660.888888888876,
23517.33093093093,
13278.112713472483,
8481.49619047619,
5050.9047619047615,
3949.275613275612,
2822.4996947496966,
2168.4787157287165,
1748.8686813186823]
```

```
mydata_one = mydata[['Age','Spending Score (1-100)']]
data=[]
for n in range(1,10):
    kmeans = (KMeans(n_clusters = n ,init='k-means++', n_init =
10 ,max_iter=400,
                    tol=0.0001, random_state= 45 ) )
    kmeans.fit(mydata_one)
    data.append(kmeans.inertia_)

plt.figure(1 , figsize = (15 ,6))
plt.plot(np.arange(1 , 10) , data , '*' )
plt.plot(np.arange(1 , 10) , data , '-' , alpha = 0.5)
plt.xlabel('Number of Clusters') , plt.ylabel('trails')
plt.show()
```




```
kmeans = KMeans(n_clusters = 4, random_state = 45)
k_fit = kmeans.fit(mydata_one)
clusters = k_fit.labels_

plt.figure(figsize = [9, 7], clear = False)
clusters = k_fit.labels_
```

```

centers = k_fit.cluster_centers_
plt.scatter(mydata_one['Age'],mydata_one['Spending Score (1-100)'],c =
clusters,s = 30,cmap = "viridis")
plt.scatter(centers[:, 0],centers[:, 1],c = "red",s = 100,alpha =
0.8);

```



```

data = pd.DataFrame({"CustomerID": mydata['CustomerID'], "Group":
(k_fit.labels_ + 1)})
data.head(10)

```

	CustomerID	Group
0	1	2
1	2	3
2	3	4
3	4	3
4	5	2
5	6	3
6	7	4
7	8	3
8	9	4
9	10	3

11. Split the data into dependent and independent variables

```
# independent variable
```

```
X = mydata.iloc[:,0:4]
```

```
X.head()
```

	CustomerID	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	0.019231	15	39
1	2	0.057692	15	81
2	3	0.038462	16	6
3	4	0.096154	16	77
4	5	0.250000	17	40

```
# dependent variable
```

```
y = mydata.iloc[:,4:]
```

```
y.head()
```

	Income Cluster
0	3
1	3
2	3
3	3
4	3

12. Split the data into training and testing

```
X_train, X_test, y_train, y_test =
```

```
train_test_split(X,y,test_size=0.3,random_state=1)
```

```
X_train.head()
```

	CustomerID	Age	Annual Income (k\$)	Spending Score (1-100)
116	117	0.865385	65	43
67	68	0.961538	48	48
78	79	0.096154	54	52
42	43	0.576923	39	36
17	18	0.038462	21	66

```
X_test.head()
```

	CustomerID	Age	Annual Income (k\$)	Spending Score (1-100)
58	59	0.173077	46	51
40	41	0.903846	38	35
34	35	0.596154	33	14
102	103	0.942308	62	59
184	185	0.442308	99	39

```
y_train.head()
```

	Income Cluster
116	4
67	1
78	4
42	1
17	3


```
y_test.head()
```

	Income Cluster
58	1
40	1
34	1
102	4
184	5

```
X_train
```

	CustomerID	Age	Annual Income (k\$)	Spending Score (1-100)
116	117	0.865385	65	43
67	68	0.961538	48	48
78	79	0.096154	54	52
42	43	0.576923	39	36
17	18	0.038462	21	66
...
133	134	0.250000	72	71
137	138	0.269231	73	73
72	73	0.807692	50	49
140	141	0.750000	75	5
37	38	0.230769	34	73

```
[140 rows x 4 columns]
```

```
X_test
```

	CustomerID	Age	Annual Income (k\$)	Spending Score (1-100)
58	59	0.173077	46	51
40	41	0.903846	38	35
34	35	0.596154	33	14
102	103	0.942308	62	59
184	185	0.442308	99	39
198	199	0.269231	137	18
95	96	0.115385	60	52
4	5	0.250000	17	40
29	30	0.096154	29	87
168	169	0.346154	87	27
171	172	0.192308	87	75
18	19	0.653846	23	29
11	12	0.326923	19	99
89	90	0.615385	58	46
110	111	0.903846	63	52
118	119	0.634615	67	43
159	160	0.230769	78	73
35	36	0.057692	33	81
136	137	0.500000	73	7
59	60	0.673077	46	46
51	52	0.288462	42	60
16	17	0.326923	21	35

44	45	0.596154	39	28
94	95	0.269231	60	42
31	32	0.057692	30	73
162	163	0.019231	81	5
38	39	0.346154	37	26
28	29	0.423077	29	31
193	194	0.384615	113	91
27	28	0.326923	28	61
47	48	0.173077	40	47
165	166	0.346154	85	75
194	195	0.557692	120	16
177	178	0.173077	88	69
176	177	0.769231	88	15
97	98	0.173077	60	50
174	175	0.653846	88	13
73	74	0.807692	50	56
69	70	0.269231	48	47
172	173	0.346154	87	10
108	109	0.961538	63	43
107	108	0.692308	63	46
189	190	0.346154	103	85
14	15	0.365385	20	13
56	57	0.634615	44	50
19	20	0.326923	23	98
114	115	0.000000	65	48
39	40	0.038462	37	75
185	186	0.230769	99	97
124	125	0.096154	70	29
98	99	0.576923	61	42
123	124	0.403846	69	91
119	120	0.615385	67	57
53	54	0.788462	43	60
33	34	0.000000	33	92
179	180	0.326923	93	90
181	182	0.269231	97	86
106	107	0.923077	63	50
199	200	0.230769	137	83
138	139	0.019231	74	10

Y_train

134	0
66	1
26	3
113	4
168	5
..	
67	1
192	2
117	4
47	1

```
172      5
Name: Income Cluster, Length: 160, dtype: int32
```

13. Build the Model

```
# classification algorithm
from sklearn.svm import SVC
classifier_model = SVC(decision_function_shape='ovo')
```

14. Train the model

```
classifier_model.fit(X_train,y_train.values.flatten())
SVC(decision_function_shape='ovo')
```

15. Test the model

```
pred_y = classifier_model.predict(X_test)
pred_y[0:5]
array([1, 1, 3, 4, 5], dtype=int32)
```

16. Measure the performance using Evaluation Metrics

```
print('Classification Report: ')
from sklearn.metrics import
classification_report,accuracy_score,f1_score,hamming_loss,confusion_m
atrix,roc_auc_score
```

```
print(classification_report(y_test, pred_y))
```

```
Classification Report:
              precision    recall  f1-score   support


     0           0.86       1.00       0.92         6
     1           1.00       0.87       0.93        15
     2           0.00       0.00       0.00         4
     3           0.91       1.00       0.95        10
     4           0.93       1.00       0.96        13
     5           0.73       0.92       0.81        12

 accuracy                   0.88         60
 macro avg           0.74       0.80       0.76         60
 weighted avg       0.84       0.88       0.85         60
```

```
print('Confusion Matrix: ')
sns.heatmap(confusion_matrix(y_test,pred_y))
```

Confusion Matrix:

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f9df91125d0>
```



```
print('F1 Score: ',f1_score(y_test,pred_y, average='weighted'))
```

F1 Score: 0.8547856464523131

Hamming loss gives the fraction of labels that are incorrectly predicted

```
print('Hamming Loss: ',hamming_loss(y_test,pred_y))
```

Hamming Loss: 0.11666666666666667

```
print('Accuracy: ',accuracy_score(y_test,pred_y))
```

Accuracy: 0.8833333333333333